

# STAT545HOMEWORK3

Xiangyu Xu

September 21, 2017

## Problem1

```
if(!require('ggplot2')){
install.packages("ggplot2")
library(ggplot2)
}

## Loading required package: ggplot2

load_mnist <- function() {
  filename <- "C:/Users/Joey/Desktop/STAT545HW3"

  load_image_file <- function(filename) {

    ret = list()

    f = file(filename, 'rb')

    readBin(f, 'integer', n=1, size=4, endian='big')

    ret$n = readBin(f, 'integer', n=1, size=4, endian='big')

    nrow = readBin(f, 'integer', n=1, size=4, endian='big')

    ncol = readBin(f, 'integer', n=1, size=4, endian='big')

    x = readBin(f, 'integer', n=ret$n*nrow*ncol, size=1, signed=F)

    ret$x = matrix(x, ncol=nrow*ncol, byrow=T)

    close(f)

    ret

  }

  load_label_file <- function(filename) {

    f = file(filename, 'rb')

    readBin(f, 'integer', n=1, size=4, endian='big')
```

```

n = readBin(f, 'integer', n=1, size=4, endian='big')

y = readBin(f, 'integer', n=n, size=1, signed=F)

close(f)

y

}

train <- load_image_file('C:/Users/Joey/Desktop/STAT545HW3/train-
images.idx3-ubyte')

train$y <- load_label_file('C:/Users/Joey/Desktop/STAT545HW3/train-
labels.idx1-ubyte')
}

show_digit <- function(arr784, col=gray(12:1/12), ...) {

  image(matrix(arr784, nrow=28)[,28:1], col=col, ...)

}

load_mnist()
digits=train$x[1:1000,]
labels=train$y[1:1000]

my_kmeans <- function(digits, K, N){
  #K means number of clusters
  #N means number of running K-means
  terminal_loss <- rep(0,N)
  n <- dim(digits)[1]
  f <- dim(digits)[2]
  for (ini in 1:N){
    set.seed(ini)
    #randomly assign observations to clusters
    No_of_cluster <- ceiling(runif(n,0,1) * K)
    last_No_of_cluster <- rep(0,n)
    mean_cluster <- matrix(0, K, f)
    eu_distance <- matrix(0, n, K)
    loss <- c(0)
    loop_loss <- 0
    while(!identical(No_of_cluster,last_No_of_cluster)){
      loop_loss <- loop_loss + 1
      for (k in 1:K){
        index <- which(No_of_cluster == k)
        if(sum(index) > 0){

```

```

    #calculate the corresponding cluster center for every cluster
    mean_cluster[k,] <- apply(digits[index, ], 2, mean)
  }
  #if empty cluster happens
  else{
    mean_cluster[k,] <- rnorm(f,0,1)
  }
}
for(k in 1:k){
  #calculate the euclidean distance
  eu_distance[,k] <- apply(t((t(digits) - mean_cluster[k,])^2), 1, sum)
  last_No_of_cluster <- No_of_cluster
}
No_of_cluster <- apply(eu_distance, 1, which.min)
loss[loop_loss] <- sum(apply(eu_distance, 1, min))
}
if(ini == 1){
  Final_loss_seq <- loss[-1]
  Final_Cluster_Assignment <- No_of_cluster
  Final_Cluster_Parameter <- mean_cluster
  terminal_loss[ini] <- loss[length(loss)]
}
else{
  if(loss[length(loss)]<Final_loss_seq[length(Final_loss_seq)]){
    Final_loss_seq <- loss[-1]
    Final_Cluster_Assignment <- No_of_cluster
    Final_Cluster_Parameter <- mean_cluster
  }
  else{
    Final_loss_seq <- Final_loss_seq
    Final_Cluster_Assignment <- No_of_cluster
    Final_Cluster_Parameter <- mean_cluster
  }
  terminal_loss[ini] <- loss[length(loss)]
}
}
return(list(Final_Cluster_Parameter <- Final_Cluster_Parameter,
            Final_Cluster_Assignment <- Final_Cluster_Assignment,
            Final_loss_seq <- Final_loss_seq,
            terminal_loss <- terminal_loss))
}

```

3

When the cluster assignment to every vector stop changing, the iteration can be stopped.



4. k=5



```


k_5 <- my_kmeans(digits,5,25)
par(mfcol <- c(2,3))

```

```
## NULL  
for (k in 1:5){  
  show_digit(k_5[[1]][k,])  
}
```

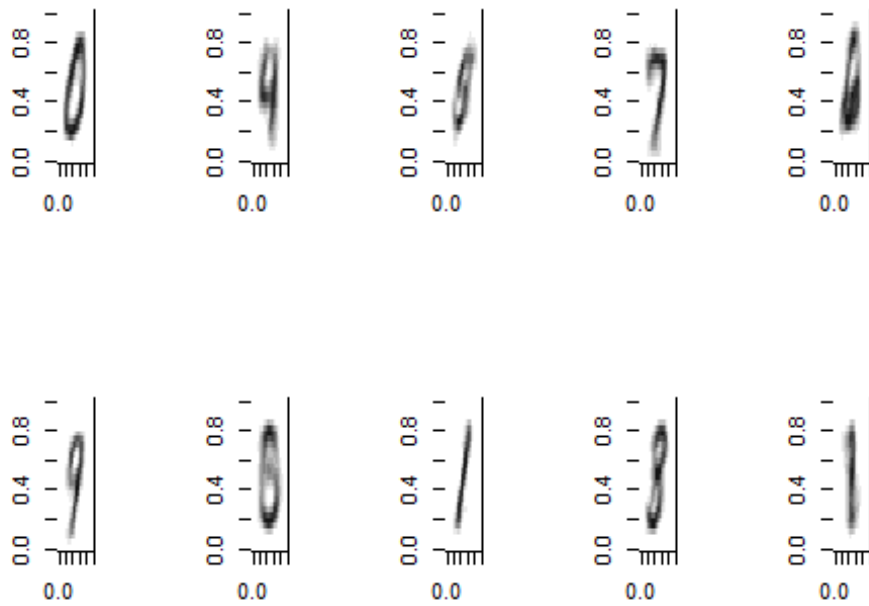






k=10

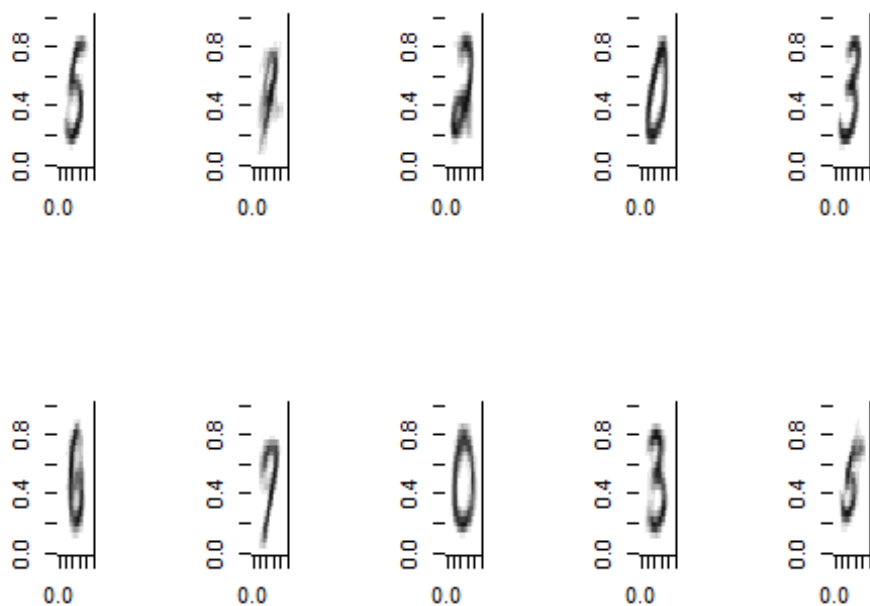
```
k_10=my_kmeans(digits,10,20)
par(mfcol=c(2,5))
for (k in 1:10){
  show_digit(k_10[[1]][k,])
}
```



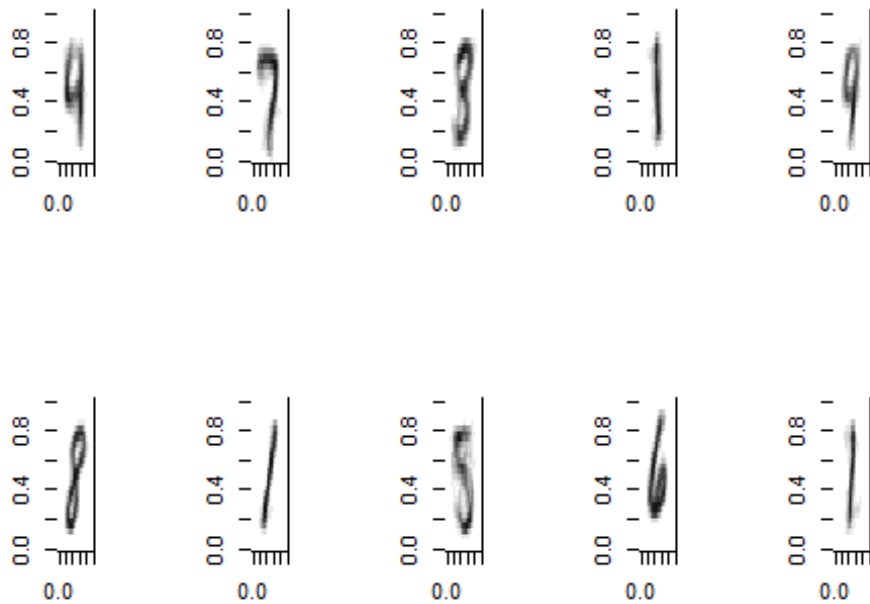
k=20

```
k_20=my_kmeans(digits,20,20)
par(mfcol=c(2,5))
for (k in 1:10){
  show_digit(k_20[[1]][k,])
}
```





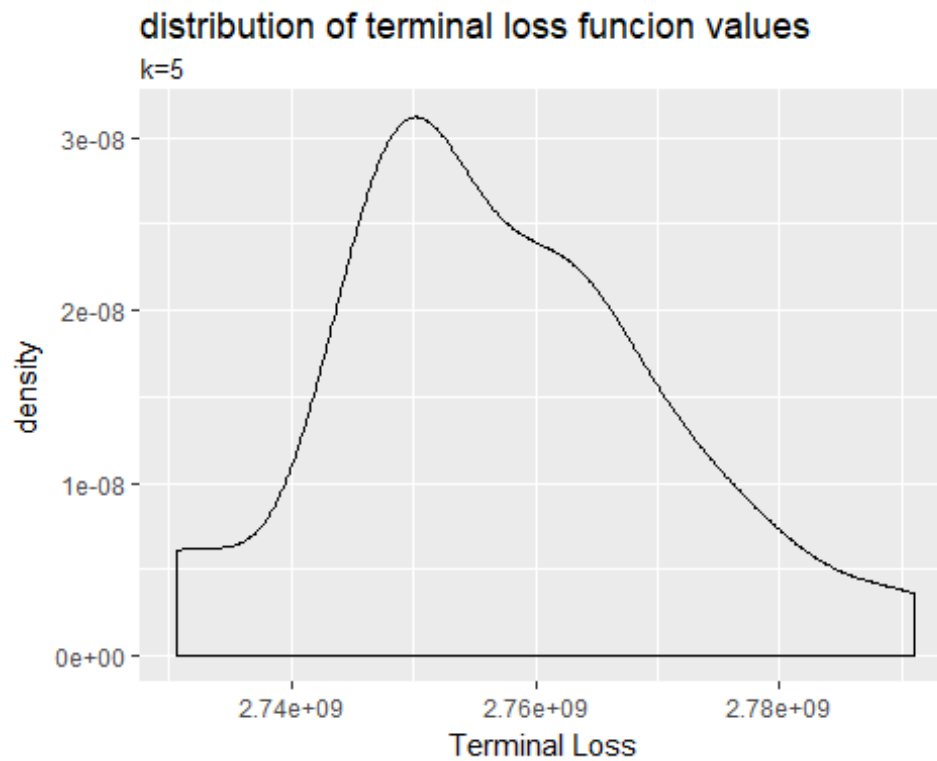
```
par(mfcol=c(2,5))
for (k in 11:20){
  show_digit(k_20[[1]][k,])
}
```



5. k=5

```
k_5_df <- data.frame(k_5[[4]])
ggplot(k_5_df, aes(x <- data.frame(k_5[[4]]))) + geom_density(colour="black")
+ ggtitle(title <- "distribution of terminal loss function values", subtitle
<- "k=5" ) + labs(x="Terminal Loss")
```


## Don't know how to automatically pick scale for object of type data.frame.  
Defaulting to continuous.



k=10

```
k_10_df <- data.frame(k_10[[4]])
ggplot(k_10_df, aes(x <- data.frame(k_10[[4]]))) +
  geom_density(colour="black") + ggtitle(title <- "distribution of terminal
loss function values", subtitle <- "k=10" ) + labs(x="Terminal Loss")


## Don't know how to automatically pick scale for object of type data.frame.
Defaulting to continuous.
```



k=20

```
k_20_df <- data.frame(k_20[[4]])
ggplot(k_20_df, aes(x <- data.frame(k_20[[4]]))) +
  geom_density(colour="black") + ggtitle(title <- "distribution of terminal
loss function values", subtitle <- "k=20" ) + labs(x="Terminal Loss")

## Don't know how to automatically pick scale for object of type data.frame.
Defaulting to continuous.
```



Apparently, if we choose  $k=1$ , the error would be the largest. As the  $k$  increases, error will decrease, and it can become 0 if  $k$  equal to the number of data points. So we should strike a balance between maximum compression of data using a single cluster, and maximum accuracy by assigning each data point to its own cluster. We could implement the elbow method. Choose a number of  $k$  so that adding another cluster doesn't give much better modeling of the data. More precisely, if we plot the percentage of variance against the  $k$ , the first cluster will add much information, but at the elbow point, the gain of variance explained will drop, thus the corresponding  $k$  to the elbow point is the value we should choose.

Bonus 7.

```
get_prototype <- function(data){
  dim_data <- dim(data)[1]
  for (ii in 1:dim_data){
    distance_medoids <- sum((t(data) - data[ii,])^2)
    if(ii == 1){
      final_distance_medoids <- distance_medoids
      center <- ii
    }
    else{
      if(distance_medoids <= final_distance_medoids){
        final_distance_medoids <- distance_medoids
        center <- ii
      }
      else{

```

```

        final_distance_medoids <- final_distance_medoids
        center <- center
    }
}
}
return(data[center,])
}

my_kmedoids <- function(digits, K, N){
  #K means number of clusters
  #N means number of running K-means
  terminal_loss <- rep(0,N)
  n <- dim(digits)[1]
  f <- dim(digits)[2]
  for (ini in 1:N){
    set.seed(ini)
    #randomly assign observations to clusters
    No_of_cluster <- ceiling(runif(n,0,1) * K)
    last_No_of_cluster <- rep(0,n)
    mean_cluster <- matrix(0, K, f)
    eu_distance <- matrix(0, n, K)
    loss <- c(0)
    loop_loss <- 0
    while(!identical(No_of_cluster,last_No_of_cluster)){
      loop_loss <- loop_loss + 1
      for (k in 1:K){
        index <- which(No_of_cluster == k)
        if(sum(index) > 0){
          #calculate the corresponding cluster center for every cluster
          mean_cluster[k,] <- get_prototype(digits[index,])
        }
      }
      for(k in 1:k){
        #calculate the euclidean distance
        eu_distance[,k] <- apply(t((t(digits) - mean_cluster[k,])^2), 1, sum)
        last_No_of_cluster <- No_of_cluster
      }
      No_of_cluster <- apply(eu_distance, 1, which.min)
      loss[loop_loss] <- sum(apply(eu_distance, 1, min))
    }
    if(ini == 1){
      Final_loss_seq <- loss[-1]
      Final_Cluster_Assignment <- No_of_cluster
      Final_Cluster_Parameter <- mean_cluster
      terminal_loss[ini] <- loss[length(loss)]
    }
    else{
      if(loss[length(loss)]<Final_loss_seq[length(Final_loss_seq)]){
        Final_loss_seq <- loss[-1]
        Final_Cluster_Assignment <- No_of_cluster
      }
    }
  }
}

```

```

        Final_Cluster_Parameter <- mean_cluster
    }
    else{
        Final_loss_seq <- Final_loss_seq
        Final_Cluster_Assignment <- No_of_cluster
        Final_Cluster_Parameter <- mean_cluster
    }
    terminal_loss[ini] <- loss[length(loss)]
}
}
return(list(Final_Cluster_Parameter <- Final_Cluster_Parameter,
            Final_Cluster_Assignment <- Final_Cluster_Assignment,
            Final_loss_seq <- Final_loss_seq,
            terminal_loss <- terminal_loss))
}

```

k=5



```

bonus_k_5 <- my_kmedoids(digits,5,20)
par(mfcol <- c(2,3))



## NULL


for (k in 1:5){
show_digit(bonus_k_5[[1]][k,])
}

```



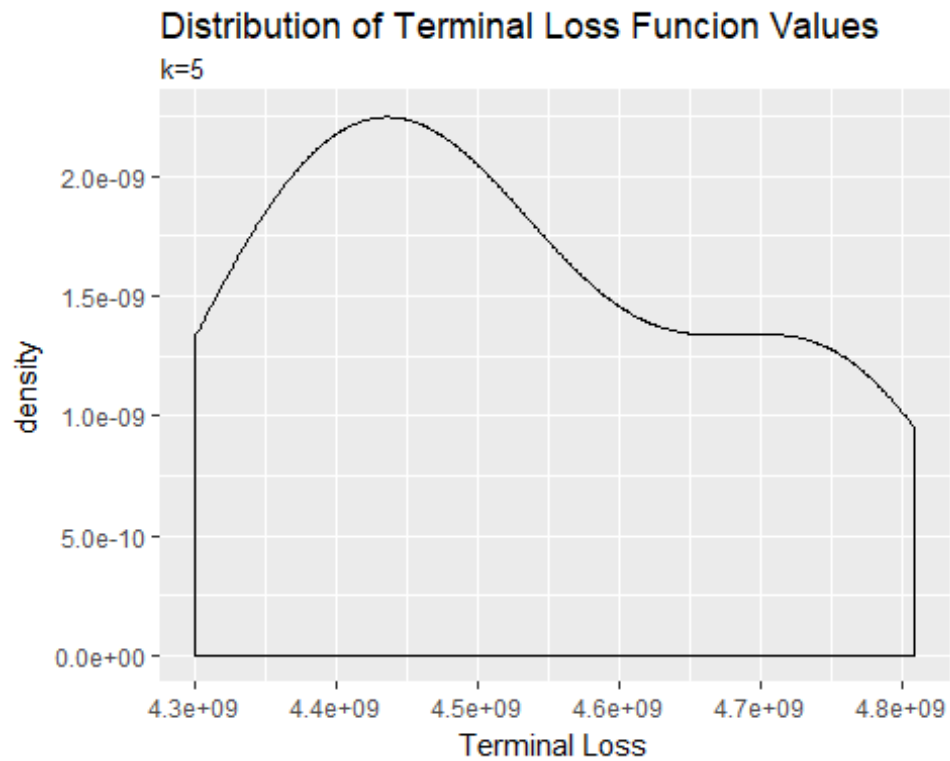






```
bonus_k_5_df <- data.frame(bonus_k_5[[4]])
ggplot(bonus_k_5_df, aes(x <- data.frame(bonus_k_5[[4]]))) +
  geom_density(colour="black") + ggtitle(title <- "Distribution of Terminal
Loss Funcion Values", subtitle <- "k=5" ) + labs(x="Terminal Loss")

## Don't know how to automatically pick scale for object of type data.frame.
Defaulting to continuous.
```







k=10



```
bonus_k_10 <- my_kmedoids(digits,10,20)
par(mfcol <- c(2,5))



## NULL



for (k in 1:10){
  show_digit(bonus_k_10[[1]][k,])
}
```












```
bonus_k_10_df <- data.frame(bonus_k_10[[4]])
ggplot(bonus_k_10_df, aes(x <- data.frame(bonus_k_10[[4]]))) +
  geom_density(colour="black") + ggtitle(title <- "Distribution of Terminal
Loss Function Values", subtitle <- "k=10" ) + labs(x="Terminal Loss")
```







```
## Don't know how to automatically pick scale for object of type data.frame.  
Defaulting to continuous.
```







k=20



```
bonus_k_20 <- my_kmedoids(digits,20,20)  
par(mfcol <- c(2,5))  
  
## NULL  
  
for (k in 1:10){  
  show_digit(bonus_k_20[[1]][k,])  
}
```















```
par(mfcol <- c(2,5))
```



```
## NULL
```



```
for (k in 11:20){  
  show_digit(bonus_k_20[[1]][k,])  
}
```
















```
bonus_k_20_df <- data.frame(bonus_k_20[[4]])
ggplot(bonus_k_20_df, aes(x = data.frame(bonus_k_20[[4]]))) +
  geom_density(colour="black") + ggtitle(title <- "Distribution of Terminal
Loss Function Values", subtitle <- "k=20" ) + labs(x="Terminal Loss")
```

```
## Don't know how to automatically pick scale for object of type data.frame.
Defaulting to continuous.
```



Compared with K-means, K-medoids has less variance on terminal loss, which indicates k-medoids would be less sensitive to the effect caused by initialization. However, the terminal loss of k-medoids is higher than k-means, that means k-means do a better job on the effect of clustering.

Problem2.

```
knitr::include_graphics('c:/Users/Joey/Desktop/stat545hw3-1.png')
```

$$= \pi_{s_1}^1 \prod_{t=1}^{T-1} P_Y(Y_t = y_t \mid S_t = s_t) A_{s_t, s_{t+1}} P_Y(Y_T = y_T \mid S_T = s_T)$$

No, this is not a probability vector.

No, they are not probability vectors.

```
knitr::include_graphics('c:/Users/Joey/Desktop/stat545hw3-3.png')
```

$$P(S_t = i \mid Y) = \frac{\alpha_i^t \beta_i^t}{\alpha^t * \beta^t}$$

```
knitr::include_graphics('c:/Users/Joey/Desktop/stat545hw3-4.png')
```

$$P(S_t = i, S_{t+1} = j | Y) = \frac{\alpha_i^t \beta_i^t}{\alpha^t * \beta^t} A_{ij}$$

```
knitr::include_graphics('c:/Users/Joey/Desktop/stat545hw3-5.png')
```

$$\alpha^t = (\alpha^{t-1})^T \text{diag}(A) B^t$$

```
knitr::include_graphics('c:/Users/Joey/Desktop/stat545hw3-6.png')
```

$$\beta^t = A(\text{diag}(\beta^{t+1}) B^t)$$

```
knitr::include_graphics('c:/Users/Joey/Desktop/stat545hw3-7.png')
```

For  $\alpha^t$ , we start from  $\alpha^1 = \text{diag}(\pi^1) B^1$  and then solve  $\alpha^2, \alpha^3, \dots, \alpha^T$ .

For  $\beta^t$ , we start from  $\beta^t = B^T$ , then we solve  $\beta^{T-1}, \beta^{T-2}, \dots, \beta^1$

```
knitr::include_graphics('c:/Users/Joey/Desktop/stat545hw3-8.png')
```

$$\underline{O}(N^2(T - t)) + O(N^2t) = O(N^2T)$$

Because the cost of computation would be pretty high.